

Group 1A Appendix

```
library(tidyverse)
library(visdat)
library(kableExtra)
library(broom)
library(ggrepel)
```

```
# Load the data
stock <- read_csv("Data/SampleA.csv")
market <- read_csv("Data/Market.csv")
```

```
# Transform to Date format
stock <- stock |>
  mutate(
    year = str_extract(Date, "\\d{4}"),
    month = str_extract(Date, "(?<=M)\\d+"),
    month = str_pad(month, width = 2, pad = "0"),
    Date = paste0(year, "-", month),
    Date = as.Date(paste0(Date, "-01"))
  ) |> select(-year, -month)
market <- market |>
  mutate(
    year = str_extract(Date, "\\d{4}"),
    month = str_extract(Date, "(?<=M)\\d+"),
    month = str_pad(month, width = 2, pad = "0"),
    Date = paste0(year, "-", month),
    Date = as.Date(paste0(Date, "-01"))
  ) |> select(-year, -month)
```

```
##/ tbl-cap: "Industry codes"
industry_codes <- tibble(Code = c("B", "C", "D", "E", "F", "G", "H", "I"),
  Industry = c("Mining", "Construction", "Manufacturing",
```

```
"Transportation and Public Utilities","Wholesale Trade",
"Retail Trade","Finance, Insurance and Real Estate","Services"))
```

```
##/ label: tbl-ind_sum
##/ tbl-cap: "Industry summary statistics"
# Extract industry from stock names
long <- stock |>
  pivot_longer(-Date, names_to = "stock", values_to = "ret")
stock_ind <- long |>
  mutate(ind = str_extract(stock, "[A-Za-z]+"))
industry_prop <- stock_ind |>
  distinct(stock, ind) |>
  count(ind, name = "n_stocks") |>
  mutate(prop = n_stocks / sum(n_stocks)) |>
  arrange(desc(prop))
ind_summary <- industry_codes |>
  left_join(industry_prop, by = c("Code" = "ind")) |>
  select(Code, Industry, n_stocks, prop)
kable(ind_summary)
```

```
#Check duplicate value
stock %>% filter(duplicated(.))
```

```
# Check for missing months
seq_months <- tibble(Date = seq(min(stock$Date, na.rm = TRUE),
                                max(stock$Date, na.rm = TRUE),
                                by = "month"))
missing_months <- seq_months |>
  anti_join(stock |> distinct(Date), by = "Date")
```

```
# Compute z-scores for each stock
stock_z <- stock |> mutate(across(-Date, ~ scale(.)[, 1], .names = "{.col}_z"))
stock_z_long <- stock_z |>
  pivot_longer(
    cols = ends_with("_z"),
    names_to = "Stock_z",
    values_to = "Z_Score"
  ) |> mutate(Stock = str_remove(Stock_z, "_z"))
# Filter rows where abs(z-score) > 3 (3-sigma outliers)
outlier_df <- stock_z_long |>
  filter(abs(Z_Score) > 3) |>
```

```

select(Date, Stock, Z_Score) |>
  arrange(desc(abs(Z_Score)))
# join with raw returns
outlier_df <- outlier_df |>
  left_join(stock |>
    pivot_longer(-Date, names_to = "Stock", values_to = "Return"),
    by = c("Date", "Stock"))

```

```

##/ label: tbl-outlier_table
##/ tbl-cap: "Top 6 outliers detected using Z-scores"
top_outlier <- head(outlier_df, 6)
kable(top_outlier)

```

```

# Per-stock summary stats
summ <- long |>
  group_by(stock) |>
  summarise(
    mean = mean(ret, na.rm = TRUE),
    sd = sd(ret, na.rm = TRUE),
    min = min(ret, na.rm = TRUE),
    p25 = quantile(ret, 0.25, na.rm = TRUE),
    median = median(ret, na.rm = TRUE),
    p75 = quantile(ret, 0.75, na.rm = TRUE),
    max = max(ret, na.rm = TRUE),
    .groups = "drop"
  ) |> arrange(desc(sd))

```

```

##/ label: tbl-statistics-table
##/ tbl-cap: "Summary statistics for stock returns"
sum_table <- summ |> slice_head(n = 9)
kable(sum_table)

```

```

# Boxplot of top volatile stocks
top9 <- summ |> slice_max(sd, n = pmin(9, nrow(summ))) |> pull(stock)

```

```

##/ label: fig-boxplot
##/ fig-cap: "Return distributions of top 9 volatile stocks"
stock |>
  select(Date, all_of(top9)) |>
  pivot_longer(-Date, names_to = "stock", values_to = "ret") |>
  ggplot(aes(stock, ret)) +

```

```
geom_boxplot(outlier.alpha = 0.5) +
coord_flip() +
labs(title = "Return distributions - top9 volatile stock",
      x = NULL,
      y = "Monthly return") +
theme_minimal()
```

Principal Component Analysis (PCA)

```
##/ label: fig-scree
##/ fig-cap: "Scree plot of PCA"
# Prepare data for PCA
stock_pca <- stock |>
  select(-Date) |>
  as.matrix()
stock_pca_std <- scale(stock_pca)
# PCA
pca <- prcomp(stock_pca_std, center = FALSE, scale. = FALSE)
screeplot(pca, type = "lines")
```

```
##/ label: tbl-pca-summary
##/ tbl-cap: "Variance explained by PC1-PC3"
var_explained <- pca$sdev^2
prop_var <- var_explained / sum(var_explained)
cum_var <- cumsum(prop_var)
# Combine into a table
pc_summary <- data.frame(
  PC = paste0("PC", 1:length(var_explained)),
  Variance = round(var_explained, 4),
  Proportion = round(prop_var, 4),
  Cumulative = round(cum_var, 4))
kable(pc_summary[1:3, ])
```

```
##/ label: tbl-cor
##/ tbl-cap: "Correlation between industry mean returns and market return"
# Industry movement
ind_move <- stock_ind |>
  group_by(Date, ind) |>
  summarise(mean_ret = mean(ret))
```

```

ind_wide <- ind_move |>
  left_join(market) |>
  select(Date, ind, mean_ret, MarketReturn) |>
  pivot_wider(names_from = ind, values_from = mean_ret)
ind_wide_num <- ind_wide |>
  mutate(across(where(is.character), as.numeric)) |>
  as.data.frame() |>
  select(-Date)
# Compute correlation
cor_mat <- cor(ind_wide_num, use = "pairwise.complete.obs")
# Turn into table
cor_tbl <- as.data.frame(round(cor_mat, 3))
knitr::kable(cor_tbl)

```

```

##/ label: tbl-loadings
##/ tbl-cap: "Industry Loadings on PC1, PC2, and PC3"
load <- as.data.frame(pca$rotation[,1:3]) |>
  rownames_to_column("stock")
load$industry <- substr(load$stock, 1, 1)
industry_centroids <- load |>
  group_by(industry) |>
  summarise(PC1 = mean(PC1), PC2 = mean(PC2), PC3 = mean(PC3))
kable(industry_centroids)

```

```

##/ label: fig-pca_cor
##/ fig-cap: "Relationships of PC1-PC3 with market return"
scores <- as.data.frame(pca$x[, 1:3]) |>
  cbind(Date = stock$Date, Market = market$MarketReturn)
# Standardize
scores_std <- scores |>
  mutate(across(-Date, ~ as.numeric(scale(.))))
scores_long <- scores_std |>
  pivot_longer(cols = c(PC1, PC2, PC3), names_to = "PC", values_to = "Score")
pca_cor_labels <- scores_long |>
  group_by(PC) |>
  summarize(cor = cor(Market, Score, use = "complete.obs"),
    .groups = "drop") |>
  mutate(label = paste0("cor = ", sprintf("%.2f", cor)),
    x = Inf, y = Inf)
ggplot(scores_long, aes(x = Market, y = Score)) +
  geom_point(alpha = 0.6) +
  geom_smooth(method = "lm") +

```

```

geom_text(data = pca_cor_labels, aes(x = x, y = y, label = label),
          hjust = 1.1, vjust = 1.5, inherit.aes = FALSE) +
facet_wrap(~PC, ncol = 1, scales = "free_y") +
theme_minimal() +
labs(title = "Relationship between PC1-PC3 and market return",
      x = "Market Return", y = "PC Score")

```

Factor Modelling

```

##/ label: fig-scrree-eigen
##/ fig-cap: "Scree plot of eigenvalues for factor analysis"
### Scree plot of eigenvalues
stock_only <- stock |> select(-Date)
X <- as.matrix(stock_only)
eig_vals <- eigen(cor(X))$values
eig_df <- data.frame(PC = 1:length(eig_vals), Eigenvalue = eig_vals)
ggplot(eig_df, aes(x = PC, y = Eigenvalue)) +
  geom_line() +
  geom_point() +
  geom_hline(yintercept = 1, linetype = "dashed", color = "red") +
  theme_minimal() +
  labs(title = "Scree Plot of Eigenvalues",
        x = "Principal Component", y = "Eigenvalue")

```

```

# Prepare matrix
stock_only <- stock |> select(-Date)
X <- as.matrix(stock_only)
# Estimate 3-factor model with Promax rotation
fa <- factanal(X, factors = 3, rotation = "promax",
               scores = "Bartlett", lower = 0.05)
print(fa, digits = 3, cutoff = 0.3)

```

```

##/ label: tbl-fa-industry
##/ tbl-cap: "Average Factor Loadings and Uniqueness by Industry"
# Estimate 3-factor model with Promax rotation
fa <- factanal(X, factors = 3, rotation = "promax",
               scores = "Bartlett", lower = 0.05)
fa_tbl <- as_tibble(unclass(fa$loadings), rownames = "stock") |>
  rename(F1 = Factor1, F2 = Factor2, F3 = Factor3) |>

```

```

mutate(
  industry = substr(stock, 1, 1),
  uniqueness = as.numeric(fa$uniquenesses))
# Summarise by industry
fa_industry <- fa_tbl |>
  group_by(industry) |>
  summarise(
    mean_F1 = mean(F1, na.rm = TRUE),
    mean_F2 = mean(F2, na.rm = TRUE),
    mean_F3 = mean(F3, na.rm = TRUE),
    mean_uniqueness = mean(uniqueness, na.rm = TRUE),
    n_stocks = n(),
    .groups = "drop")
kable(fa_industry, digits = 3)

```

```

##/ label: tbl-fa-stock
##/ tbl-cap: "Top 10 Most Systematic Stocks (Lowest Uniqueness)"
fa_top10_systematic <- fa_tbl |>
  arrange(uniqueness) |>
  slice_head(n = 10) |>
  select(industry, stock, starts_with("F"), uniqueness)
knitr::kable(fa_top10_systematic, digits = 3)

```

```

##/ label: fig-fa-stock
##/ fig-cap: "Uniqueness vs strongest factor loading"
fa_stock_plot <- fa_tbl |>
  rowwise() |>
  mutate(max_loading = max(abs(c_across(dplyr::starts_with("F"))))) |> ungroup()
ggplot(fa_stock_plot, aes(x = uniqueness, y = max_loading, color = industry, label = stock))
  geom_point() +
  geom_text_repel(size = 2, max.overlaps = 100) +
  labs(
    title = "Uniqueness (idiosyncratic) vs strongest factor loading (systematic)",
    x = "Uniqueness (lower = more systematic)",
    y = "Max |Factor Loading|"
  ) + theme_minimal()

```

```

##/ label: fig-fa-loadings
##/ fig-cap: "Factor Loadings Plot"
fa_df <- tidy(fa)
ggplot(fa_df, aes(x = fl1, y = fl2, label = variable)) +

```

```

geom_segment(aes(xend = fl1, yend = fl2, x = 0, y = 0),
             arrow = arrow(length = unit(2, "mm")), linewidth = 0.2) +
geom_point(size = 0.8) +
geom_text_repel(color='red',size = 2, max.overlaps = 50) +
coord_equal() +
labs(title = "Factor Loadings Plot (F1 vs F2)",
     x = "Factor 1 Loadings", y = "Factor 2 Loadings")

```

```

##/ label: tbl-rank
##/ tbl-cap: "Top 5 Most Systematic Stocks with FA and PCA Loadings"
# Factor Analysis loadings + uniqueness
fa_tbl <- as_tibble(unclass(fa$loadings), rownames = "stock") |>
  rename(F1 = Factor1, F2 = Factor2, F3 = Factor3)
uniq_tbl <- tibble(stock = names(fa$uniquenesses),
                  uniqueness = as.numeric(fa$uniquenesses))
fa_tbl <- fa_tbl |>
  left_join(uniq_tbl, by = "stock")
# PCA loadings (rotation)
pc_tbl <- as.data.frame(pca$rotation[,1:3]) |>
  rownames_to_column("stock") |>
  rename(PC1 = PC1, PC2 = PC2, PC3 = PC3)
# Combine both
rank_tbl <- fa_tbl |>
  left_join(pc_tbl, by = "stock") |>
  mutate(industry = substr(stock, 1, 1)) |>
  arrange(uniqueness)
# Preview top 5
rank_tbl_pretty <- rank_tbl |>
  mutate(across(c(F1,F2,F3,PC1,PC2,PC3,uniqueness), ~round(.x,3))) |>
  select(industry, stock, uniqueness, F1, F2, F3, PC1, PC2, PC3)
knitr::kable(rank_tbl_pretty |> slice_head(n = 5))

```

```

##/ label: fig-fa-top5
##/ fig-cap: "Uniqueness vs strongest factor loading, highlighting top 5 picks"
# Get absolute strongest factor loading for each stock
df <- rank_tbl %>%
  rowwise() %>%
  mutate(max_loading = max(abs(c(F1, F2, F3)))) %>%
  ungroup()
# Mark top 5 recommended stocks
top5 <- c("H90000", "H88215", "H75157", "H89050", "H77466")
df <- df %>% mutate(top_pick = ifelse(stock %in% top5, "Yes", "No"))

```



```
# Plot
ggplot(df,aes(x=uniqueness,y = max_loading,color=top_pick,label = stock)) +
  geom_point(size = 3) +
  geom_text_repel(aes(label = ifelse(top_pick == "Yes", stock, "")),
    size = 3, box.padding = 0.25, max.overlaps = Inf) +
  scale_color_manual(values = c("Yes" = "red", "No" = "grey")) +
  theme_minimal() +
  labs(title = "Uniqueness vs Strongest Factor Loading",
    subtitle = "Top 5 market-like stocks highlighted in red",
    x = "Uniqueness (idiosyncratic risk)",
    y = "Strongest Factor Loading (systematic risk)",
    color = "Top 5 Pick")
```

Other

```
## Check missing values and data types
vis_miss(stock) +
  theme(axis.text.x = element_blank(),
    axis.ticks.x = element_blank())
vis_dat(stock) +
  theme(axis.text.x = element_blank(),
    axis.ticks.x = element_blank())
```

```
##/ fig-cap: "Stock Price Movement by industry"
# Time series
ggplot(ind_move, aes(x = Date)) +
  geom_line(aes(y = mean_ret), color = "black") +
  facet_wrap(~ind) + theme_minimal() +
  labs(title = "Industry mean returns over time",x = "Date",y = "Return") +
  theme(legend.position = "bottom")
```

```
# ----- Build X (T x N) from 'stock' and standardise -----
X <- stock %>%
  dplyr::select(-Date) %>%
  as.matrix()
# mean-impute each column (transparent + simple)
X <- apply(X, 2, function(x) { x[is.na(x)] <- mean(x, na.rm = TRUE); x })
# standardise over time so Euclidean distance ~ correlation structure
Y <- scale(X) # T x N
```

```

Yc <- t(Y) # N x T (rows = stocks)
# preserve stock names for later plots
if (is.null(colnames(X))) colnames(X) <- paste0("S", seq_len(ncol(X)))
rownames(Yc) <- colnames(X)

library(MASS) # isoMDS
library(mclust) # Mclust, adjustedRandIndex
library(factoextra)
# distances between stocks
d_stocks <- dist(Yc, method = "euclidean")
# Ward.D2
hc <- hclust(d_stocks, method = "ward.D2")
k <- 3
cl_hier <- cutree(hc, k = k)
# visual check
fviz_dend(hc, cex = 0, xlab = "", ylab = "", main = "", ggtheme = theme_bw(),
          lwd = 1.5, k = 3, rect = TRUE, color_labels_by_k = FALSE) +
  theme(text = element_text(size = 26), axis.text.x = element_blank(),
        plot.margin = margin(50, 20, 60, 20)) + coord_cartesian(clip = "off")
# Base on the Dendrogram, 3 cluster solution is suggested, even it is not as
# stable as 2 cluster solution, the size of clusters would be more balanced.

```

```

hc_result <- cutree(hc, k = 3) %>%
  as.data.frame() %>%
  rownames_to_column("stock") %>%
  mutate(industry = substr(stock, 1, 1)) %>%
  rename(cluster = ".")
# Count per cluster and industry, calculate percentage
cluster_industry_percent <- hc_result %>%
  group_by(cluster, industry) %>%
  summarise(count = n(), .groups = "drop") %>%
  group_by(cluster) %>%
  mutate(percent = round(100 * count / sum(count), 2)) %>%
  ungroup() %>%
  dplyr::select(-count) %>%
  pivot_wider(names_from = cluster, values_from = percent, names_prefix = "Cluster_")
knitr::kable(cluster_industry_percent)
# Cluster 1 is dominated by industry H, D, and I, cluster 2 is dominated by
# industry H, and E, and cluster 3 only contain cluster H.

```